

MIDTERM PROJECT

Course: Mining Massive Datasets

Duration: 03 weeks

I. Formation

- The midterm project is conducted in groups of **04 – 05** students.
- The student group fulfills the requirements and submits the work according to the detailed instructions below.

II. Requirements

Given the **data.csv** file contains the customer's purchase data. In which, the first line contains the title (header), the remaining lines are the corresponding data rows.

- **Member_number**: member number
- **Date**: date of purchase in “dd/mm/yyyy”
- **itemDescription**: one product name
- **year**: year of purchase
- **month**: month of purchase
- **day**: day of purchase
- **day_of_week**: day of week

For example,

Member_number	Date	itemDescription	year	month	day	day_of_week
1249	01/01/2014	citrus fruit	2014	1	1	2
1249	01/01/2014	coffee	2014	1	1	2
1381	01/01/2014	curd	2014	1	1	2
1381	01/01/2014	soda	2014	1	1	2
1440	01/01/2014	other vegetables	2014	1	1	2
1440	01/01/2014	yogurt	2014	1	1	2
1659	01/01/2014	specialty chocolate	2014	1	1	2
1659	01/01/2014	frozen vegetables	2014	1	1	2
1789	01/01/2014	hamburger meat	2014	1	1	2
1789	01/01/2014	candles	2014	1	1	2

Sample data rows in data.csv (Google Colab)

a) Task 1: Item Counting

Use **textFile()** method of **SparkContext** to read **data.csv** file and count the number of items bought by each customer for each **day**, **month**, and **year**. Results are saved 03 folders named **counters_day**, **counters_month**, and **counters_year** respectively using **saveAsTextFile()**. Result files include:

- For day: *Member_number, Day, Quantity*.
- For month: *Member_number, Month, Quantity*.
- For year: *Member_number, Year, Quantity*.

Find out the maximum number of items in a basket sold by day, then save the results to the **max_day** folder. The result content includes *Date, Maximum number*

Notices:

- Use only **RDD** functions to find results (do not use **DataFrame**).
- The result does not contain the header line, the values on each line are separated by a “,” and do not contain a “,” at the end of the line.
- Students read the result file with the **DataFrame** of **SQL Context** and display it with the **show()** function to demonstrate in the exercise..

Criteria	Score
<i>Number of items for each day</i>	0.5 point(s)
<i>Number of items for each month</i>	0.5 point(s)
<i>Number of items for each year</i>	0.5 point(s)
<i>Maximum number of items in a basket for each day</i>	0.5 point(s)
<i>Total</i>	2.0 point(s)

b) Task 2: Baskets

Use **textFile()** method of **SparkContext** to read **data.csv** and find out *the list of items bought by each customer for each day*. Results are saved to the **baskets** directory using the **saveAsTextFile()** method. The result content includes: *Member_number, Purchase date, List of items*.

Notices:

- Columns are separated by “;” and items are separated by “,”.
- Each item in a basket appears no more than 1 time.

For example,

part-00000 ×

```
1 Member_number;Date;itemDescription
2 1249;01/01/2014;citrus fruit,coffee
3 1381;01/01/2014;curd,soda
4 1440;01/01/2014;other vegetables,yogurt
5 1659;01/01/2014;specialty chocolate,frozen vegetables
6 1789;01/01/2014;hamburger meat,candles
7 1922;01/01/2014;tropical fruit,other vegetables
8 2226;01/01/2014;sausage,bottled water
9 2237;01/01/2014;bottled water,Instant food products
10 2351;01/01/2014;cleaner,shopping bags
```

List of items bought by each customer for each day

The first line is the header. Items are separated by “;”

Items in a basket are separated by “,”

Students read the result file with **SQL Context DataFrame** and display it with **show()** function to demonstrate..

Member_number	Date	itemDescription
1249	01/01/2014	citrus fruit,coffee
1381	01/01/2014	curd,soda
1440	01/01/2014	other vegetables,...
1659	01/01/2014	specialty chocola...
1789	01/01/2014	hamburger meat,ca...

An example result of Task 2

Write a program (2A) that allows the user to enter a member number and a date in the form *dd/mm/yyyy*, and then print the corresponding basket.

Write a program (2B) that allows the user to enter a member number and a positive integer *n*, and then print a result table containing the list of baskets that person has purchased in descending order of purchase date. Take only up to *n* lines of data.

Criteria	Score
Find out and save the list of baskets	1.0 point(s)
Program (2A)	0.5 point(s)
Program (2B)	0.5 point(s)
Total	2.0 point(s)

c) Task 3: Frequent Itemsets

- Read the result file in b) into a **DataFrame** named *dfBaskets*, convert the **itemDescription** column into an array of strings and display it to the screen with the **show()** method.
- Rename columns of the result data frame as *Member_number*, *Date*, and *Items*.

For example

Member_number	Date	Items
1249	01/01/2014	[coffee, citrus f...
1381	01/01/2014	[soda, curd]
1440	01/01/2014	[yogurt, other ve...
1659	01/01/2014	[specialty chocol...
1789	01/01/2014	[hamburger meat, ...]
1922	01/01/2014	[other vegetables...
2226	01/01/2014	[sausage, bottled...

An example data frame

- Use `pyspark.ml.fpm.FPGrowth` to build up a model find frequent item sets with **minSupport=1%** and association rules with **minConfidence=10%**. Students display the list of frequent item sets and association rules to the screen using **show()** method.

For example,

items	freq	antecedent	consequent	confidence
[beef]	508	[other vegetables]	[whole milk]	0.12151067323481117
[sugar]	265	[yogurt]	[whole milk]	0.12996108949416343
[oil]	223	[rolls/buns]	[whole milk]	0.12697448359659783
[chocolate]	353	[soda]	[whole milk]	0.11975223675154852
[white wine]	175			
[candy]	215			
[processed cheese]	152			
[meat]	252			

Frequent Item sets

Association rules

Write a program (3A) that finds **support** for each item in the data set with a **support threshold** s entered by the user. The results are displayed on a tabular screen containing the **item name** column and the **support** column, in which the data rows are sorted descending by support.

Write a program (3B) that allows the user to enter two item names x and y . Find and print the values **confidence**($x \rightarrow y$) and **interest**($x \rightarrow y$) to the screen.

Students study and present details of `pyspark.ml.fpm.FPGrowth`.

Criteria	Score
<i>Transform dfBaskets correctly</i>	0.5 point(s)
<i>Program (3A)</i>	0.5 point(s)
<i>Program (3B)</i>	0.5 point(s)
<i>Details of FPGrowth</i>	0.5 point(s)
<i>Total</i>	2.0 point(s)

d) Task 4: Baskets-to-Vectors

- Read **data.csv** as a **DataFrame** named **dfMembers** with two columns
 - Member_number*: Member number
 - Items*: list of items, separated by “;”, which contains all items bought by a customer and every item is unique.
- Students can display the results on the screen using **show()** function to demonstrate.

For example

```
+-----+-----+
|Member_number|          Items|
+-----+-----+
|          1249|citrus fruit,coff...|
|          1381|curd,soda,coffee,...|
|          1440|other vegetables,...|
|          1659|specialty chocola...|
```

dfMembers DataFrame

- Find the item list in the entire data set, save it in the **items** variable, where the elements are unique and sorted in ascending alphabetical order. Then, create a dictionary named

dictItems containing pairs *<item name>:<list index>*. Print the values of the two variables above to the screen.

For example

```
[ 'Instant food products', 'UHT-milk', 'abrasive cleaner', 'artif.
{ 'Instant food products': 0, 'UHT-milk': 1, 'abrasive cleaner': 2,
```

Examples of *items* and *dictItems*

- Implement the function **basket2vector(member, basket, dictItems)** taking a member number (**member**), a list of items (**basket**) separated by “,”, and **dictItems** above. The function returns a **Vectors.sparse (pyspark.ml.linalg)**, in form of multi-hot coding, with the same length as **dictItems**, elements are 0.0 or 1.0 corresponding to existing items in **basket** (for example there is item *it* in **basket**, then **dictItems[it] = 1.0**).

For example: invoke **basket2vector** and pass the first row of **dfMembers** in

```
print(basket2vector(dfMembers.first()[ 'Member_number' ],
                    dfMembers.first()[ 'Items' ],
                    dictItems))
```

```
(167, [11,30,34,61,138], [1.0,1.0,1.0,1.0,1.0])
```

Returned value from **basket2vector** for the first row of **dfMembers**

Students note that **DataFrame dfMembers** is also used in e).

Criteria	Score
Find out the item list	0.5 point(s)
Create the dictionary correctly	
Function basket2vector	0.5 point(s)
Create dfMember correctly	
Total	1.0 point(s)

e) Task 5: Similar baskets

- Create a model **MinHashLSH** in module **pyspark.ml.feature** to find out customers with similar shopping habit. Use the data frame **dfMembers** from Task 4, in which
 - inputCol** is “Items”

- *outputCol* is “Hashes”
- Number of hash tables is 08
- Train the model with *dfMembers*, apply *transform()* function in *dfMembers*, and then display the result data frame on the screen.
- Note do not change the value of *dfMembers* because it is used in later tasks.

For example

Member_number	Items	Hashes
1249	(167, [11, 30, 34, 61...]	[[2.85001106E8], ...]
1381	(167, [1, 10, 11, 28, ...]	[[3.9022841E7], [...]
1440	(167, [28, 64, 102, 1...]	[[3.41446049E8], ...]
1659	(167, [12, 14, 26, 34...]	[[1.02200615E8], ...]

Result *DataFrame* transformed from *dfMembers*

- Use **approxSimilarityJoin()** to find out pairs of customers with similar shopping habit, where **JaccardDistance** is not over **0.5**. Display the result data frame with columns including member numbers and **JaccardDistance**, in which selecting rows with positive *JaccardDistance*.

For example,

idA	idB	JaccardDistance
3124	1063	0.25
1643	4535	0.19999999999999996
1860	3605	0.25
4342	1056	0.2857142857142857
2911	3714	0.25
3715	4805	0.25

Result pairs of customers

- Use **approxNearestNeighbors()** to find out **05** customers whose shopping habit is the most similar to the one whose *member_number* is input by users. Display the result on the screen.

For example,

Member_number	Items	Hashes	distCol
1249	(167, [11, 30, 34, 61...]	[[2.85001106E8], ...]	0.0
1321	(167, [11, 30, 138], ...]	[[2.85001106E8], ...]	0.4
1263	(167, [11, 30, 61, 10...]	[[2.85001106E8], ...]	0.5
1794	(167, [11, 30, 138, 1...]	[[2.85001106E8], ...]	0.5714285714285714
4327	(167, [30, 34, 63, 76...]	[[3.43690326E8], ...]	0.5714285714285714

Top 05 customers whose shopping habit is the most similar to the first one in dfMembers

Students study and present details of **MinHashLSH** in PySpark.

Criteria	Score
<i>Use MinHashLSH to transform data correctly</i>	0.5 point(s)
<i>Find out pairs of customers with similar shopping habit</i>	0.5 point(s)
<i>Find out top 05 most similar customers to input one</i>	0.5 point(s)
<i>Study and present details of MinHashLSH</i>	0.5 point(s)
<i>Total</i>	2.0 point(s)

f) Task 7 (1.0 point): Presentation

- Student groups compose a presentation to report your work.
- **THERE IS NO PRESENTATION TEMPLATES. STUDENTS ARRANGE CONTENTS IN A LOGICAL LAYOUT BY YOURSELVES.**
- The presentation must include below contents
 - Student list: Student ID, Full name, Email, Assigned tasks, Complete percentage.
 - Briefly present approaches to solve tasks, should make use of pseudo code/diagrams.
 - AVOID EMBEDDING RAW SOURCE CODE IN THE PRESENTATION.
 - Study topics are introduced briefly with practical examples.
 - Advantages versus disadvantages
 - A table of complete percentages for each task.
 - References are presented in IEEE format.

- **Format requirements:** slide ratio of 4x3, avoid using dark background/colorful shapes because of projector quality, students ensure contents are clear enough when printing the presentation in grayscale.
- Presentation duration is **10 minutes**.

III. Submission Instructions

- Create a folder whose name is as
 <Student ID 1>_< Student ID 2>_< Student ID 3>_< Student ID 4>
- Content:
 - **source.ipynb** → source code (remain all cell outputs)
 - **source.pdf** → pdf of the notebook
 - **presentation.pdf** → presentation.
- Compress the folder to a zip file and submit by the deadline.

IV. Policy

- **Student groups submitting late get 0.0 points for each member.**
- **Wrong student IDs in the submission filename cause 0.0 points for the corresponding students.**
- **Missing required materials in the submission loses at least 50% points of the presentation.**
- **Copying source code on the internet/other students, sharing your work with other groups, etc. cause 0.0 points for all related groups.**
- **If there exist any signs of illegal copying or sharing of the assignment, then extra interviews are conducted to verify student groups' work.**

-- THE END --